

welcome: [please sign in](#)
location: [id3v2.4.0-structure](#)

Informal standard
Document: id3v2.4.0-structure.txt

M. Nilsson
1st November 2000

ID3 tag version 2.4.0 - Main Structure

Status of this document

This document is an informal standard and replaces the ID3v2.3.0 standard [ID3v2]. A formal standard will use another revision number even if the content is identical to document. The contents in this document may change for clarifications but never for added or altered functionality.

Distribution of this document is unlimited.

Abstract

This document describes the main structure of ID3v2.4.0, which is a revised version of the ID3v2 informal standard [ID3v2] version 2.3.0. The ID3v2 offers a flexible way of storing audio meta information within the audio file itself. The information may be technical information, such as equalisation curves, as well as title, performer, copyright etc.

ID3v2.4.0 is meant to be as close as possible to ID3v2.3.0 in order to allow for implementations to be revised as easily as possible.

1. Table of contents

- Status of this document
- Abstract
- 1. Table of contents
- 2. Conventions in this document
- 2. Standard overview
- 3. ID3v2 overview
 - 3.1. ID3v2 header
 - 3.2. ID3v2 extended header
 - 3.3. Padding
 - 3.4. ID3v2 footer
- 4. ID3v2 frames overview
 - 4.1. Frame header flags
 - 4.1.1. Frame status flags
 - 4.1.2. Frame format flags
- 5. Tag location
- 6. Unsynchronisation
 - 6.1. The unsynchronisation scheme

- 6.2. Synchsafe integers
- 7. Copyright
- 8. References
- 9. Author's Address

2. Conventions in this document

Text within `"` is a text string exactly as it appears in a tag. Numbers preceded with `$` are hexadecimal and numbers preceded with `%` are binary. `$xx` is used to indicate a byte with unknown content. `%x` is used to indicate a bit with unknown content. The most significant bit (MSB) of a byte is called 'bit 7' and the least significant bit (LSB) is called 'bit 0'.

A tag is the whole tag described in this document. A frame is a block of information in the tag. The tag consists of a header, frames and optional padding. A field is a piece of information; one value, a string etc. A numeric string is a string that consists of the characters "0123456789" only.

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in RFC 2119 [KEYWORDS].

3. ID3v2 overview

ID3v2 is a general tagging format for audio, which makes it possible to store meta data about the audio inside the audio file itself. The ID3 tag described in this document is mainly targeted at files encoded with MPEG-1/2 layer I, MPEG-1/2 layer II, MPEG-1/2 layer III and MPEG-2.5, but may work with other types of encoded audio or as a stand alone format for audio meta data.

ID3v2 is designed to be as flexible and expandable as possible to meet new meta information needs that might arise. To achieve that ID3v2 is constructed as a container for several information blocks, called frames, whose format need not be known to the software that encounters them. At the start of every frame is an unique and predefined identifier, a size descriptor that allows software to skip unknown frames and a flags field. The flags describes encoding details and if the frame should remain in the tag, should it be unknown to the software, if the file is altered.

The bitorder in ID3v2 is most significant bit first (MSB). The byteorder in multibyte numbers is most significant byte first (e.g. \$12345678 would be encoded \$12 34 56 78), also known as big endian and network byte order.

Overall tag structure:

```
+-----+
|      Header (10 bytes)      |
+-----+
|      Extended Header        |
+-----+
```

```
| (variable length, OPTIONAL) |  
+-----+  
|   Frames (variable length)   |  
+-----+  
|           Padding            |  
| (variable length, OPTIONAL) |  
+-----+  
| Footer (10 bytes, OPTIONAL) |  
+-----+
```

In general, padding and footer are mutually exclusive. See details in sections 3.3, 3.4 and 5.

3.1. ID3v2 header

The first part of the ID3v2 tag is the 10 byte tag header, laid out as follows:

ID3v2/file identifier	"ID3"
ID3v2 version	\$04 00
ID3v2 flags	%abcd0000
ID3v2 size	4 * %0xxxxxxx

The first three bytes of the tag are always "ID3", to indicate that this is an ID3v2 tag, directly followed by the two version bytes. The first byte of ID3v2 version is its major version, while the second byte is its revision number. In this case this is ID3v2.4.0. All revisions are backwards compatible while major versions are not. If software with ID3v2.4.0 and below support should encounter version five or higher it should simply ignore the whole tag. Version or revision will never be \$FF.

The version is followed by the ID3v2 flags field, of which currently four flags are used.

a - Unsynchronisation

Bit 7 in the 'ID3v2 flags' indicates whether or not unsynchronisation is applied on all frames (see section 6.1 for details); a set bit indicates usage.

b - Extended header

The second bit (bit 6) indicates whether or not the header is followed by an extended header. The extended header is described in section 3.2. A set bit indicates the presence of an extended header.

c - Experimental indicator

The third bit (bit 5) is used as an 'experimental indicator'. This flag SHALL always be set when the tag is in an experimental stage.

d - Footer present

Bit 4 indicates that a footer (section 3.4) is present at the very end of the tag. A set bit indicates the presence of a footer.

All the other flags MUST be cleared. If one of these undefined flags are set, the tag might not be readable for a parser that does not know the flags function.

The ID3v2 tag size is stored as a 32 bit synchsafe integer (section 6.2), making a total of 28 effective bits (representing up to 256MB).

The ID3v2 tag size is the sum of the byte length of the extended header, the padding and the frames after unsynchronisation. If a footer is present this equals to ('total size' - 20) bytes, otherwise ('total size' - 10) bytes.

An ID3v2 tag can be detected with the following pattern:

\$49 44 33 yy yy xx zz zz zz zz

Where yy is less than \$FF, xx is the 'flags' byte and zz is less than \$80.

3.2. Extended header

The extended header contains information that can provide further insight in the structure of the tag, but is not vital to the correct parsing of the tag information; hence the extended header is optional.

Extended header size	4 * %0xxxxxxx
Number of flag bytes	\$01
Extended Flags	\$xx

Where the 'Extended header size' is the size of the whole extended header, stored as a 32 bit synchsafe integer. An extended header can thus never have a size of fewer than six bytes.

The extended flags field, with its size described by 'number of flag bytes', is defined as:

%0bcd0000

Each flag that is set in the extended header has data attached, which comes in the order in which the flags are encountered (i.e. the data for flag 'b' comes before the data for flag 'c'). Unset flags cannot have any attached data. All unknown flags MUST be unset and their corresponding data removed when a tag is modified.

Every set flag's data starts with a length byte, which contains a value between 0 and 128 (\$00 - \$7f), followed by data that has the field length indicated by the length byte. If a flag has no attached data, the value \$00 is used as length byte.

b - Tag is an update

If this flag is set, the present tag is an update of a tag found earlier in the present file or stream. If frames defined as unique are found in the present tag, they are to override any corresponding ones found in the earlier tag. This flag has no corresponding data.

Flag data length \$00

c - CRC data present

If this flag is set, a CRC-32 [ISO-3309] data is included in the extended header. The CRC is calculated on all the data between the header and footer as indicated by the header's tag length field, minus the extended header. Note that this includes the padding (if there is any), but excludes the footer. The CRC-32 is stored as an 35 bit synchsafe integer, leaving the upper four bits always zeroed.

Flag data length \$05
Total frame CRC 5 * %0xxxxxxx

d - Tag restrictions

For some applications it might be desired to restrict a tag in more ways than imposed by the ID3v2 specification. Note that the presence of these restrictions does not affect how the tag is decoded, merely how it was restricted before encoding. If this flag is set the tag is restricted as follows:

Flag data length \$01
Restrictions %ppqrrstt

p - Tag size restrictions

- 00 No more than 128 frames and 1 MB total tag size.
- 01 No more than 64 frames and 128 KB total tag size.
- 10 No more than 32 frames and 40 KB total tag size.
- 11 No more than 32 frames and 4 KB total tag size.

q - Text encoding restrictions

- 0 No restrictions
- 1 Strings are only encoded with ISO-8859-1 [ISO-8859-1] or UTF-8 [UTF-8].

r - Text fields size restrictions

- 00 No restrictions
- 01 No string is longer than 1024 characters.
- 10 No string is longer than 128 characters.
- 11 No string is longer than 30 characters.

Note that nothing is said about how many bytes is used to represent those characters, since it is encoding dependent. If a text frame consists of more than one string, the sum of the strings is restricted as stated.

s - Image encoding restrictions

- 0 No restrictions
- 1 Images are encoded only with PNG [PNG] or JPEG [JFIF].

t - Image size restrictions

- 00 No restrictions
- 01 All images are 256x256 pixels or smaller.
- 10 All images are 64x64 pixels or smaller.
- 11 All images are exactly 64x64 pixels, unless required otherwise.

3.3. Padding

It is OPTIONAL to include padding after the final frame (at the end of the ID3 tag), making the size of all the frames together smaller than the size given in the tag header. A possible purpose of this padding is to allow for adding a few additional frames or enlarge existing frames within the tag without having to rewrite the entire file. The value of the padding bytes must be \$00. A tag MUST NOT have any padding between the frames or between the tag header and the frames. Furthermore it MUST NOT have any padding when a tag footer is added to the tag.

3.4. ID3v2 footer

To speed up the process of locating an ID3v2 tag when searching from the end of a file, a footer can be added to the tag. It is REQUIRED to add a footer to an appended tag, i.e. a tag located after all audio data. The footer is a copy of the header, but with a different identifier.

ID3v2 identifier	"3DI"
ID3v2 version	\$04 00
ID3v2 flags	%abcd0000
ID3v2 size	4 * %0xxxxxxx

4. ID3v2 frame overview

All ID3v2 frames consists of one frame header followed by one or more fields containing the actual information. The header is always 10 bytes and laid out as follows:

Frame ID	\$xx xx xx xx (four characters)
Size	4 * %0xxxxxxx
Flags	\$xx xx

The frame ID is made out of the characters capital A-Z and 0-9. Identifiers beginning with "X", "Y" and "Z" are for experimental frames and free for everyone to use, without the need to set the experimental bit in the tag header. Bear in mind that someone else might have used the same identifier as you. All other identifiers are either used or reserved for future use.

The frame ID is followed by a size descriptor containing the size of the data in the final frame, after encryption, compression and unsynchronisation. The size is excluding the frame header ('total frame size' - 10 bytes) and stored as a 32 bit synchsafe integer.

In the frame header the size descriptor is followed by two flag bytes. These flags are described in section 4.1.

There is no fixed order of the frames' appearance in the tag, although it is desired that the frames are arranged in order of significance concerning the recognition of the file. An example of such order: UFID, TIT2, MCDI, TRCK ...

A tag **MUST** contain at least one frame. A frame must be at least 1 byte big, excluding the header.

If nothing else is said, strings, including numeric strings and URLs [URL], are represented as ISO-8859-1 [ISO-8859-1] characters in the range \$20 - \$FF. Such strings are represented in frame descriptions as <text string>, or <full text string> if newlines are allowed. If nothing else is said newline character is forbidden. In ISO-8859-1 a newline is represented, when allowed, with \$0A only.

Frames that allow different types of text encoding contains a text encoding description byte. Possible encodings:

- \$00 ISO-8859-1 [ISO-8859-1]. Terminated with \$00.
- \$01 UTF-16 [UTF-16] encoded Unicode [UNICODE] with BOM. All strings in the same frame SHALL have the same byteorder. Terminated with \$00 00.
- \$02 UTF-16BE [UTF-16] encoded Unicode [UNICODE] without BOM. Terminated with \$00 00.
- \$03 UTF-8 [UTF-8] encoded Unicode [UNICODE]. Terminated with \$00.

Strings dependent on encoding are represented in frame descriptions as <text string according to encoding>, or <full text string according to encoding> if newlines are allowed. Any empty strings of type \$01 which are NULL-terminated may have the Unicode BOM followed by a Unicode NULL (\$FF FE 00 00 or \$FE FF 00 00).

The timestamp fields are based on a subset of ISO 8601. When being as precise as possible the format of a time string is yyyy-MM-ddTHH:mm:ss (year, "-", month, "-", day, "T", hour (out of 24), ":", minutes, ":", seconds), but the precision may be reduced by removing as many time indicators as wanted. Hence valid timestamps are

yyyy, yyyy-MM, yyyy-MM-dd, yyyy-MM-ddTHH, yyyy-MM-ddTHH:mm and yyyy-MM-ddTHH:mm:ss. All time stamps are UTC. For durations, use the slash character as described in 8601, and for multiple non-

contiguous dates, use multiple strings, if allowed by the frame definition.

The three byte language field, present in several frames, is used to describe the language of the frame's content, according to ISO-639-2 [ISO-639-2]. The language should be represented in lower case. If the language is not known the string "XXX" should be used.

All URLs [URL] MAY be relative, e.g. "picture.png", "../doc.txt".

If a frame is longer than it should be, e.g. having more fields than specified in this document, that indicates that additions to the frame have been made in a later version of the ID3v2 standard. This is reflected by the revision number in the header of the tag.

4.1. Frame header flags

In the frame header the size descriptor is followed by two flag bytes. All unused flags MUST be cleared. The first byte is for 'status messages' and the second byte is a format description. If an unknown flag is set in the first byte the frame MUST NOT be changed without that bit cleared. If an unknown flag is set in the second byte the frame is likely to not be readable. Some flags in the second byte indicates that extra information is added to the header. These fields of extra information is ordered as the flags that indicates them. The flags field is defined as follows (l and o left out because their resemblance to one and zero):

```
%0abc0000 %0h00kmp
```

Some frame format flags indicate that additional information fields are added to the frame. This information is added after the frame header and before the frame data in the same order as the flags that indicates them. I.e. the four bytes of decompressed size will precede the encryption method byte. These additions affects the 'frame size' field, but are not subject to encryption or compression.

The default status flags setting for a frame is, unless stated otherwise, 'preserved if tag is altered' and 'preserved if file is altered', i.e. %00000000.

4.1.1. Frame status flags

a - Tag alter preservation

This flag tells the tag parser what to do with this frame if it is unknown and the tag is altered in any way. This applies to all kinds of alterations, including adding more padding and reordering the frames.

- | | |
|---|----------------------------|
| 0 | Frame should be preserved. |
| 1 | Frame should be discarded. |

b - File alter preservation

This flag tells the tag parser what to do with this frame if it is unknown and the file, excluding the tag, is altered. This does not apply when the audio is completely replaced with other audio data.

- 0 Frame should be preserved.
- 1 Frame should be discarded.

c - Read only

This flag, if set, tells the software that the contents of this frame are intended to be read only. Changing the contents might break something, e.g. a signature. If the contents are changed, without knowledge of why the frame was flagged read only and without taking the proper means to compensate, e.g. recalculating the signature, the bit **MUST** be cleared.

4.1.2. Frame format flags**h - Grouping identity**

This flag indicates whether or not this frame belongs in a group with other frames. If set, a group identifier byte is added to the frame. Every frame with the same group identifier belongs to the same group.

- 0 Frame does not contain group information
- 1 Frame contains group information

k - Compression

This flag indicates whether or not the frame is compressed. A 'Data Length Indicator' byte **MUST** be included in the frame.

- 0 Frame is not compressed.
- 1 Frame is compressed using zlib [zlib] deflate method.
If set, this requires the 'Data Length Indicator' bit to be set as well.

m - Encryption

This flag indicates whether or not the frame is encrypted. If set, one byte indicating with which method it was encrypted will be added to the frame. See description of the ENCR frame for more information about encryption method registration. Encryption should be done after compression. Whether or not setting this flag requires the presence of a 'Data Length Indicator' depends on the specific algorithm used.

- 0 Frame is not encrypted.
- 1 Frame is encrypted.

n - Unsynchronisation

This flag indicates whether or not unsynchronisation was applied to this frame. See section 6 for details on unsynchronisation. If this flag is set all data from the end of this header to the end of this frame has been unsynchronised. Although desirable, the presence of a 'Data Length Indicator' is not made mandatory by unsynchronisation.

- 0 Frame has not been unsynchronised.
- 1 Frame has been unsynchronised.

p - Data length indicator

This flag indicates that a data length indicator has been added to the frame. The data length indicator is the value one would write as the 'Frame length' if all of the frame format flags were zeroed, represented as a 32 bit synchsafe integer.

- 0 There is no Data Length Indicator.
- 1 A data length Indicator has been added to the frame.

5. Tag location

The default location of an ID3v2 tag is prepended to the audio so that players can benefit from the information when the data is streamed. It is however possible to append the tag, or make a prepend/append combination. When deciding upon where an unembedded tag should be located, the following order of preference SHOULD be considered.

1. Prepend the tag.
2. Prepend a tag with all vital information and add a second tag at the end of the file, before tags from other tagging systems. The first tag is required to have a SEEK frame.
3. Add a tag at the end of the file, before tags from other tagging systems.

In case 2 and 3 the tag can simply be appended if no other known tags are present. The suggested method to find ID3v2 tags are:

1. Look for a prepended tag using the pattern found in section 3.1.
2. If a SEEK frame was found, use its values to guide further searching.
3. Look for a tag footer, scanning from the back of the file.

For every new tag that is found, the old tag should be discarded unless the update flag in the extended header (section 3.2) is set.

6. Unsynchronisation

The only purpose of unsynchronisation is to make the ID3v2 tag as compatible as possible with existing software and hardware. There is no use in 'unsynchronising' tags if the file is only to be processed only by ID3v2 aware software and hardware. Unsynchronisation is only useful with tags in MPEG 1/2 layer I, II and III, MPEG 2.5 and AAC files.

6.1. The unsynchronisation scheme

Whenever a false synchronisation is found within the tag, one zeroed byte is inserted after the first false synchronisation byte. The format of synchronisations that should be altered by ID3 encoders is as follows:

```
%11111111 111xxxxx
```

and should be replaced with:

```
%11111111 00000000 111xxxxx
```

This has the side effect that all \$FF 00 combinations have to be altered, so they will not be affected by the decoding process. Therefore all the \$FF 00 combinations have to be replaced with the \$FF 00 00 combination during the unsynchronisation.

To indicate usage of the unsynchronisation, the unsynchronisation flag in the frame header should be set. This bit **MUST** be set if the frame was altered by the unsynchronisation and **SHOULD NOT** be set if unaltered. If all frames in the tag are unsynchronised the unsynchronisation flag in the tag header **SHOULD** be set. It **MUST NOT** be set if the tag has a frame which is not unsynchronised.

Assume the first byte of the audio to be \$FF. The special case when the last byte of the last frame is \$FF and no padding nor footer is used will then introduce a false synchronisation. This can be solved by adding a footer, adding padding or unsynchronising the frame and add \$00 to the end of the frame data, thus adding more byte to the frame size than a normal unsynchronisation would. Although not preferred, it is allowed to apply the last method on all frames ending with \$FF.

It is preferred that the tag is either completely unsynchronised or not unsynchronised at all. A completely unsynchronised tag has no false synchronisations in it, as defined above, and does not end with \$FF. A completely non-unsynchronised tag contains no unsynchronised frames, and thus the unsynchronisation flag in the header is cleared.

Do bear in mind, that if compression or encryption is used, the unsynchronisation scheme **MUST** be applied afterwards. When decoding an unsynchronised frame, the unsynchronisation scheme **MUST** be reversed first, encryption and decompression afterwards.

6.2. Synchsafe integers

In some parts of the tag it is inconvenient to use the unsynchronisation scheme because the size of unsynchronised data is not known in advance, which is particularly problematic with size descriptors. The solution in ID3v2 is to use synchsafe integers, in which there can never be any false synchs. Synchsafe integers are integers that keep its highest bit (bit 7) zeroed, making seven bits out of eight available. Thus a 32 bit synchsafe integer can store 28 bits of information.

Example:

255 (%11111111) encoded as a 16 bit synchsafe integer is 383 (%00000001 01111111).

7. Copyright

Copyright (C) Martin Nilsson 2000. All Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that a reference to this document is included on all such copies and derivative works. However, this document itself may not be modified in any way and reissued as the original document.

The limited permissions granted above are perpetual and will not be revoked.

This document and the information contained herein is provided on an 'AS IS' basis and THE AUTHORS DISCLAIMS ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

8. References

[ID3v2] Martin Nilsson, 'ID3v2 informal standard'.

<url:http://www.id3.org/id3v2.3.0.txt>

[ISO-639-2] ISO/FDIS 639-2.

'Codes for the representation of names of languages, Part 2: Alpha-3 code.' Technical committee / subcommittee: TC 37 / SC 2

[ISO-3309] ISO 3309

'Information Processing Systems--Data Communication High-Level Data Link Control Procedure--Frame Structure', IS 3309, October 1984, 3rd Edition.

[ISO-8859-1] ISO/IEC DIS 8859-1.

'8-bit single-byte coded graphic character sets, Part 1: Latin

alphabet No. 1.' Technical committee / subcommittee: JTC 1 / SC 2

[JFIF] 'JPEG File Interchange Format, version 1.02'

<url:http://www.w3.org/Graphics/JPEG/jfif.txt>

[KEYWORDS] S. Bradner, 'Key words for use in RFCs to Indicate Requirement Levels', RFC 2119, March 1997.

<url:ftp://ftp.isi.edu/in-notes/rfc2119.txt>

[MPEG] ISO/IEC 11172-3:1993.

'Coding of moving pictures and associated audio for digital storage media at up to about 1,5 Mbit/s, Part 3: Audio.'

Technical committee / subcommittee: JTC 1 / SC 29

and

ISO/IEC 13818-3:1995

'Generic coding of moving pictures and associated audio information, Part 3: Audio.'

Technical committee / subcommittee: JTC 1 / SC 29

and

ISO/IEC DIS 13818-3

'Generic coding of moving pictures and associated audio information, Part 3: Audio (Revision of ISO/IEC 13818-3:1995)'

[PNG] 'Portable Network Graphics, version 1.0'

<url:http://www.w3.org/TR/REC-png-multi.html>

[UNICODE] The Unicode Consortium,

'The Unicode Standard Version 3.0', ISBN 0-201-61633-5.

<url:http://www.unicode.org/unicode/standard/versions/Unicode3.0.htm>

[URL] T. Berners-Lee, L. Masinter & M. McCahill, 'Uniform Resource Locators (URL)', RFC 1738, December 1994.

<url:ftp://ftp.isi.edu/in-notes/rfc1738.txt>

[UTF-8] F. Yergeau, 'UTF-8, a transformation format of ISO 10646', RFC 2279, January 1998.

<url:ftp://ftp.isi.edu/in-notes/rfc2279.txt>

[UTF-16] F. Yergeau, 'UTF-16, an encoding of ISO 10646', RFC 2781, February 2000.

<url:ftp://ftp.isi.edu/in-notes/rfc2781.txt>

[ZLIB] P. Deutsch, Aladdin Enterprises & J-L. Gailly, 'ZLIB Compressed Data Format Specification version 3.3', RFC 1950, May 1996.

<url:ftp://ftp.isi.edu/in-notes/rfc1950.txt>

9. Author's Address

Written by

Martin Nilsson
Rydsvägen 246 C. 30
SE-584 34 Linköping
Sweden

Email: [nilsson at id3.org](mailto:nilsson@id3.org)

id3v2.4.0-structure (last edited 2012-10-08 22:15:41 by localhost)

[Copyright](#) © 1998-2024 by their respective owners